




**Plasma-PEPSC Webinar  
28 May 2024**

## **alpaka Parallel Programming Library**



Hello everyone. Thank you for participating. In this presentation I will initially introduce alpaka generally. Then I will go over an alpaka program and explain it step by step. Lastly I will talk about some performance tests and the contributors of alpaka.

## alpaka – Abstraction Library for Parallel Kernel Acceleration

### alpaka is...

- A parallel programming library: Accelerate your code by exploiting your hardware's parallelism!
- An abstraction library independent of hardware ecosystem: Create portable code that runs on CPUs and GPUs!
- Free & open-source software




Alpaka is a parallel programming library. You can accelerate your code by exploiting your hardware's parallelism. It is an abstraction library independent of hardware ecosystem.

With alpaka you can create create portable code which runs on different GPUs, and on CPUs.

Alpaka is free and open-source.

And lastly it is yet another library using an animal as logo.




## Problem of HPC Systems?

## Heterogenous Hardware Ecosystem!

### TOP500

- 1 **Frontier(USA)** 1.194 Exaflop/s,  
**AMD EPYC CPU + AMD Instinct GPU**
- 2 **Aurora(USA)** 585 Petaflop/s,  
**Intel Xeon CPU + Intel GPU Max**
- 3 **Eagle(USA)** 561 Petaflop/s,  
**Intel Xeon CPU + Nvidia GPU H100**
- 4 **Fugaku(Japan)** 442 Petaflop/s,  
**Fujitsu A64FX CPU**
- 5 **Lumi(Finland)** 380 Petaflop/s,  
**AMD EPYC CPU + AMD Instinct GPU**



**CASUS**  
CENTRUM FOR ADVANCED  
COMPUTATIONAL SCIENCES

[www.top500.org](http://www.top500.org)

**THE LIST**

**11/2023 Highlights**

The 42<sup>nd</sup> edition of the TOP500 shows that nine of upgraded entries in the top 10 but the Frontier system still remains the only that exceeds machine with an HPL score of 1.1 ExaFlops.

Frontier is composed of the Oak Ridge National Laboratory Truetime, USA remains the No. 1 system as the TOP500 and is still the only system reported with exascale performance exceeding one ExaFlop. Frontier brought the year position back to the USA one year ago as the June 2023 listing and has since been reinforced with an HPL score of 1.1 ExaFlops.

Frontier is based on the latest HPE Cray EX10 architecture and is equipped with AMD EPYC CPUs and AMD Instinct MI300X GPUs. The system has 8,494,944 cores, a power efficiency rating of 33.64 gigaFlops/Watt, and relies on HPE's Singularity II network for data transfer.

Frontier system is the largest Leadership Computing Facility (LCF), USA is currently being commissioned and will be full scale exceed Frontier with a peak performance of 2 ExaFlops. It was equipped with a replacement on top of the first system achieving 100 Petaflops on the HPL benchmark which secured the No. 2 spot on the TOP500.

Aurora is built by Intel based on the HPE Cray EX - Intel Exascale Compute Blade which uses Intel Xeon CPU/Intel Xeon Phi processors and Intel Xeon GPU/Intel Xeon Phi accelerators connected through HPE's Singularity II network interconnect.

Aurora is built by Intel based on the HPE Cray EX - Intel Exascale Compute Blade which uses Intel Xeon CPU/Intel Xeon Phi processors and Intel Xeon GPU/Intel Xeon Phi accelerators connected through HPE's Singularity II network interconnect.

Frontier system installed in the Microsoft Azure cloud in the USA is easily listed as No. 3. This Microsoft M207 system is based on Intel Xeon Platinum 8480C processors with 4,096 accelerators and achieved an HPL score of 361 PetaFlops.

[read more >](#)

**List Statistics**


**Heterogeneous Systems**

1	Frontier - HPE Cray EX10, AMD Instinct MI300X, AMD EPYC 9611, AMD Instinct MI300X, Singularity II, HPL
2	Aurora - HPE Cray EX - Intel Exascale Compute Blade, Intel Xeon Phi 9700, Intel Xeon Phi 9700, Intel Xeon Phi 9700, Singularity II, HPL
3	Eagle - Microsoft M207, Intel Xeon Phi 9700, Intel Xeon Phi 9700, Intel Xeon Phi 9700, Singularity II, HPL
4	Fugaku - Fujitsu A64FX, AMD Instinct MI300X, Singularity II, HPL
5	Lumi - AMD EPYC 9611, AMD Instinct MI300X, AMD Instinct MI300X, Singularity II, HPL
6	Summit - AMD EPYC 9611, AMD Instinct MI300X, AMD Instinct MI300X, Singularity II, HPL
7	Summit - AMD EPYC 9611, AMD Instinct MI300X, AMD Instinct MI300X, Singularity II, HPL
8	Summit - AMD EPYC 9611, AMD Instinct MI300X, AMD Instinct MI300X, Singularity II, HPL
9	Summit - AMD EPYC 9611, AMD Instinct MI300X, AMD Instinct MI300X, Singularity II, HPL
10	Summit - AMD EPYC 9611, AMD Instinct MI300X, AMD Instinct MI300X, Singularity II, HPL

If we just go over the list of top HPC systems, we can see that hardware ecosystem is quite heterogenous.

The first one in the list, Frontier in USA is using AMD GPUs, Aurora is using Intel GPUs, the third one Eagle is using Nvidia GPUs, Fugaku in Japan is using CPUs which are actually CPUs with ARM architecture and Lumi in Finland is using AMD GPUs.

And these different vendors propose different programming APIs. Even they are using different colors in their logos and publications.

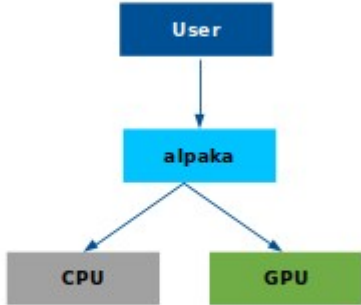


**Without alpaka**

- Hardware ecosystem is heterogenous, platforms are not inter-operable → parallel programs not easily portable


**alpaka: one API to rule them all**

- Abstraction** (not hiding!) of the underlying hardware, compiler and OS
  - No default device, built-in functions, language extensions
- Easy change of the backend in code**
- Direct usage of vendor APIs, not depend on "unified APIs"**
  - GPU Backends: Hip (AMD), Cuda (NVidia), SYCL (Intel GPUs)
  - One can use vendor profilers and debuggers (Cuda, HIP...) for alpaka code!
  - CPU Backends: OpenMp, Threads, TbbBlocks
- Zero abstraction overhead for Kernel execution!**
- Heterogenous Programming:** Using different backends in a synchronized manner.



```

graph TD
    User[User] --> alpaka[alpaka]
    alpaka --> CPU[CPU]
    alpaka --> GPU[GPU]
            
```


Alpaka in A Nutshell | 4

Hence currently HPC Platforms are not interoperable, or in other words programs are not portable.

## Alpaka provides one API to support all different GPUs and CPU backends.

**Abstraction** (but not hiding!) of the underlying hardware, compiler and OS is the main approach of Alpaka.

Alpaka does not have default device, built-in functions, language extensions, default stream like in cuda

**It is Easy to change the backend in code**

**Alpaka code directly uses vendor APIs, does not depend on “unified APIs”. Produces the same code that a vendor API would generate. Hence alpaka has Zero abstraction overhead for Kernel execution!**

Supported GPU Backends are Hip (AMD), Cuda (NVidia), SYCL (Intel GPUs)

alpaka user can use vendor profilers and debuggers (Cuda, HIP...) for his alpaka code!

Supported CPU Backends: OpenMp, Threads, TbbBlocks **Heterogenous Programming:** Using different backends in a synchronized manner

Find us on GitHub!

**alpaka library:** <https://www.github.com/alpaka-group/alpaka>

- Full source code and many examples, Issue tracker

**The documents:** <https://alpaka.readthedocs.io/en/latest/>

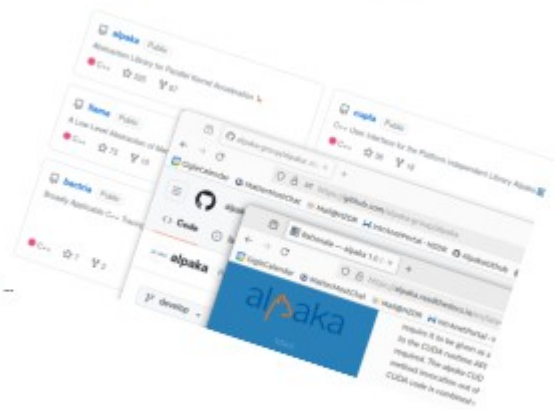
- Installation guide
- Cheatsheet
- Abstraction model and the rationale behind alpaka

**Project group:** <https://www.github.com/alpaka-group>

- Contains all alpaka-related projects, documentation, samples, ...

alpaka is a free software (MPL 2.0)

 Public License



You can Find alpaka on GitHub!

**The Github includes** Full source code and many examples, and an Issue tracker

The documents pages at [readthedocs](https://alpaka.readthedocs.io/en/latest/) includes:

- Installation guide
- Cheatsheet
- Abstraction model and the rationale behind alpaka

Alpaka Project group: <https://www.github.com/alpaka-group>

Contains all alpaka-related projects, documentation, samples, ...

Among those softwares:

- cupla easy porting from cuda,
- vikunja is for high level **algorithms**.

alpaka is a free software (MPL 2.0)

## Programming with alpaka

- C++ only!
- alpaka is written entirely in C++17. Coming soon: C++20.
- Header-only library. No additional runtime dependency.  
`#include <alpaka/alpaka.hpp>` is enough!
- Supports a wide range of modern C++ compilers (g++, clang++, Apple LLVM, MSVC)
- Portable across operating systems: Linux, macOS, Windows



Alpaka is a library for C++. and it is written entirely in C++17. In a short time we will be using C++20 features and compiling on : C++20.

It is a Header-only library. No additional runtime dependency is used. Only Compile time dependency is Boost.

Including the header file `alpaka/alpaka.hpp` to the cpp code would be enough!

Alpaka Supports a wide range of modern C++ compilers (g++, clang++, Apple LLVM, MSVC)

and it is Portable across operating systems: Linux, macOS, Windows

**alpaka**

**CASUS**  
CENTRE FOR ADVANCED SCIENTIFIC COMPUTING

## Installation and Examples

### I-Install Dependencies

- alpaka requires **Boost**, **Cmake** and a modern C++ compiler (g++, clang++, Visual C++, ...)
  - Linux: `sudo apt install libboost-all-dev` (DEB)
  - MacOS: `brew install boost` (using homebrew, <https://brew.sh>)
  - Windows: `vcpkg install boost` (using vcpkg, <https://github.com/microsoft/vcpkg>)
- Depending on your target platform you may need additional packages
  - NVIDIA GPUs: CUDA Toolkit (<https://developer.nvidia.com/cuda-toolkit>)
  - AMD GPUs: ROCm / HIP (<https://rocm.docs.amd.com/en/latest/index.html>)
  - Intel GPUs: OneAPI Toolkit (<https://www.intel.com/content/www/us/en/developer/tools/oneapi/base-toolkit.html#gs.9x3lnh>)
- CMake is the preferred system for building and installing
  - Linux: `sudo apt install cmake` (DEB). macOS and Windows: Download the installer from <https://cmake.org/download/>

Alpaka in A Notebook | 7

## Installation of Alpaka and Building Examples

Initially we need to install dependencies

alpaka requires Boost as a compile time dependency for compilation. For the configuration of the build system Cmake and for the compilation a C++ compiler (g++, clang++, Visual C++, ...) is needed.

Depending on your target platform you may need additional packages of Cuda, Rocm or Intel OneAPI toolkits.

CMake is the preferred system for building and installing

## II - Compiling and running examples

- You can build all examples at once from your build directory:
  - configure the build with setting some cmake variables according to your system  

```
cmake -Dalpaka_BUILD_EXAMPLES=ON -DCMAKE_BUILD_TYPE=Release -Dalpaka_ACC_CPU_B_SEQ_T_SEQ_ENABLE=ON -Dalpaka_ACC_GPU_CUDA_ENABLE=ON ..
```
  - build the examples  

```
cmake --build . --config Release
```
  - alpaka/build/example/ directory will include compiled examples.  
 e.g. alpaka/build/example/vectorAdd directory will include the executable `vectorAdd`
- Run all examples from the build directory of alpaka  

```
ctest example/
```
- Run all tests from the build directory of alpaka  

```
ctest test/
```
- Examples can be re-compiled and run in their corresponding directories under build directory if there is a code change in the source tree.  

```
cd alpaka/build/example/vectorAdd
cmake --build . (or run the make command if make file is there)
```

```
1 git clone https://github.com/alpaka-group/alpaka.git
2 cd alpaka
3 mkdir build
4 cd build/
5 cmake -Dalpaka_BUILD_EXAMPLES=ON -DCMAKE_BUILD_TYPE=Release -Dalpaka_ACC_CPU_B_SEQ_T_SEQ_ENABLE=ON -Dalpaka_ACC_GPU_CUDA_ENABLE=ON ..
6 cmake --build . --config Release
7 cd example/vectorAdd/
8 ./vectorAdd
```

After installing boost and cmake and compilation tools We don't need to install alpaka files to compile examples. We can just directly compile and run.

In compiling examples or any program using alpaka; Setting cmake variables are important. The user needs to configure the build with setting the cmake variables according to her/his system:

```
cmake -Dalpaka_BUILD_EXAMPLES=ON -DCMAKE_BUILD_TYPE=Release -Dalpaka_ACC_CPU_B_SEQ_T_SEQ_ENABLE=ON -Dalpaka_ACC_GPU_CUDA_ENABLE=ON
```

These backend settings doesn't mean the user has to use these backends in the code but means they are available and can be used by the user. ( // Example selects which one?)

After building, Runnin all examples by **ctest example** command or all tests by **ctest test** command is possible.



### III - Install alpaka Library

Download alpaka: `git clone -b develop https://github.com/alpaka-group/alpaka.git`

- In the terminal/powershell, switch to the downloaded alpaka directory:

```
cd /path/to/alpaka
```

- Create a build directory and switch to it:

```
mkdir build
```

```
cd build
```

- Configure build directory (If default directories is ok for you or you are planning to use alpaka from build directory; you can omit the install prefix cmake variable)

```
cmake -DCMAKE_INSTALL_PREFIX=/some/other/path/..
```

- **Install alpaka without compiling!** alpaka installation will reside in `/some/other/path/`.

```
cmake --install .
```

- You should now have a complete alpaka installation in the directory you chose earlier.

For Detailed information: <https://github.com/alpaka-group/alpaka-workshop-slides/tree/develop>

```
git clone https://github.com/alpaka-group/alpaka.git
cd alpaka
mkdir build
cd build
cmake -DCMAKE_INSTALL_PREFIX=/some/other/path/ ..
cmake --install .
```

Installation of alpaka is quite strait-forward. Compilation is not needed, just 2 steps are needed: configuring and installing.

Notice that you don't need to set cmake variables depending on our system for installation; because those variables will be just before the compilation of your project. Namely configuring your project.



## IV - Create your first alpaka project

### Option 1

1. After installation, create your cmake file as the one below. Use `find_package(alpaka REQUIRED)`
2. Create your example app code (possible under `src` directory of the directory)
3. Inside the directory of your project run commands:  

```
mkdir build
```

```
cd build
```
4. Configure the example according to your system:  

```
cmake ..
```
5. Build your code:  

```
make
```

```
##
# CMakeLists.txt for the myHelloWorld example
cmake_minimum_required(VERSION 3.22)

set(TARGET_NAME myHelloWorld)
project(${TARGET_NAME} LANGUAGES CXX)
# Find alpaka.
find_package(alpaka REQUIRED)
# Add executable.
alpaka_add_executable(
    ${TARGET_NAME}
    src/myHelloWorld.cpp
)
target_link_libraries(
    ${TARGET_NAME}
    PUBLIC alpaka::alpaka
)
```

### Option 2

1. Copy one of the examples to a folder outside the alpaka source tree. The folder will be your project directory.
2. Remove `ExampleDefaultAcc` and `getAccName` from the code. Select a backend (accelerator) inside the code instead of `ExampleDefaultAcc`.
3. Clone alpaka under your project directory
4. Add alpaka directory to `CmakeLists.txt` by `add_subdirectory`
5. Repeat steps 3-5 of **Option 1** under your project directory.

```
# CMakeLists.txt for using alpaka by adding alpaka
# as a subdirectory to your project.
cmake_minimum_required(VERSION 3.22)

# Project.
set(TARGET_NAME myVectorAddProject)
project(${TARGET_NAME} LANGUAGES CXX)

# add alpaka directory name(cloned into your project)
add_subdirectory(./alpaka)

# Add executable.
alpaka_add_executable(
    ${TARGET_NAME}
    src/vectorAdd.cpp
)
target_link_libraries(
    ${TARGET_NAME}
    PUBLIC alpaka::alpaka
)
```

```
vectorAddProject/
├── alpaka
│   ├── CMakeLists.txt
│   ├── CHANGELOG.md
│   ├── cmake
│   ├── CMakePresets.json
│   ├── CONTRIBUTING.md
│   ├── docs
│   ├── example
│   ├── include
│   ├── LICENSE
│   ├── README.md
│   ├── README_SYCL.md
│   ├── scripts
│   ├── test
│   └── thirdParty
├── build
│   ├── CMakeCache.txt
│   ├── CMakeFiles
│   ├── cmake_install.cmake
│   ├── Makefile
│   ├── myVectorAddProject
│   ├── CMakeLists.txt
│   └── src
└── src
    └── vectorAdd.cpp
```

Creating your alpaka project can be achieved in 2 ways. First way is, creating your **cmakelists** file and your code file in a directory you chose.

Secondly you can just copy one of the examples as a directory and change the code a little bit.

The important point for both cases is that. Select all possible accelerators while configuring the build tree by cmake, that means, user sets the corresponding cmake variables to activate accelerators during the configuration by cmake.

### Tenets of Thread-Parallel Programming

- **Grid, block and thread based parallelisation model.**

The model is instantiated differently on different processors, because of cache size and speed, the synchronization mechanism, or simply the CPU-GPU difference.

- **Large number of threads** should run the same code (**kernel**) on different data in parallel.
- **Indexing of threads.** Each thread should work on a different data portion or do a specific task, therefore each thread has an index accessible in kernel.

- **Extent:** A vector representing the sizes along dimensions.

In 3d an extent is {Width,Length,Height}



- **Dimensions:** Set of dimension names. {X-dimension, Y-dimension, Z-dimension}

- **Number of Dimensions**



Grid-Block Extent: Vec{3}  
Block-Thread Extent: Vec{4}



Grid-Block Extent: Vec{3} or Vec{1,3}  
Block-Thread Extent: Vec{4,5}

Before moving on a small alpaka code I would like to talk about basic concepts of thread-parallel programming.

Alpaka uses Grid-block-thread based parallelisation model. The model is instantiated differently on different processors, because of cache size and the cache speed, the synchronization mechanism, or simply the difference between CPU-GPUs. By using grid-block-thread abstraction the execution can be optimally adapted to the available hardware.

Secondly the assumption is that Large number of threads should run the same code (kernel) on different data in a parallel manner.

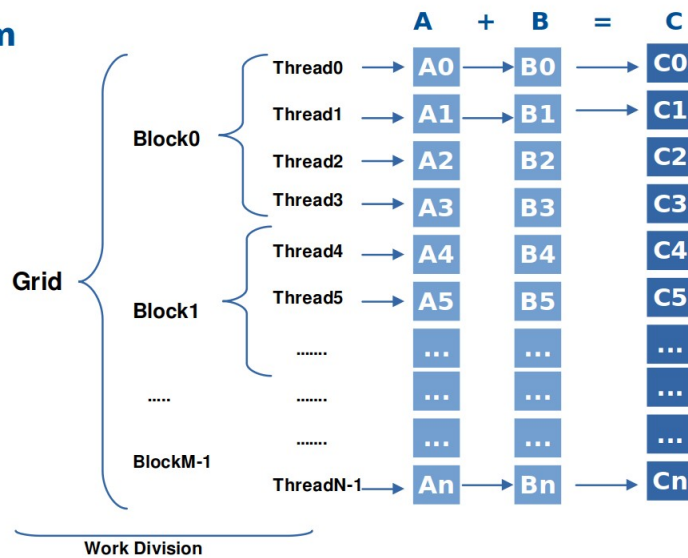
Lastly; Indexing of threads is needed. Each thread should work on a different data portion or do a specific task, therefore each thread has an index accessible in kernel.

There are 3 terms I would like to describe. The extent means the sizes along each dimensions. In 3d for example extent is a 3 item vector of {width,length and height}. The term Dimensions means “set of dimension names” although in daily english we use dimensions as the extent.

Lastly number of dimensions is the size of the extent vector.

## Vector Addition Problem

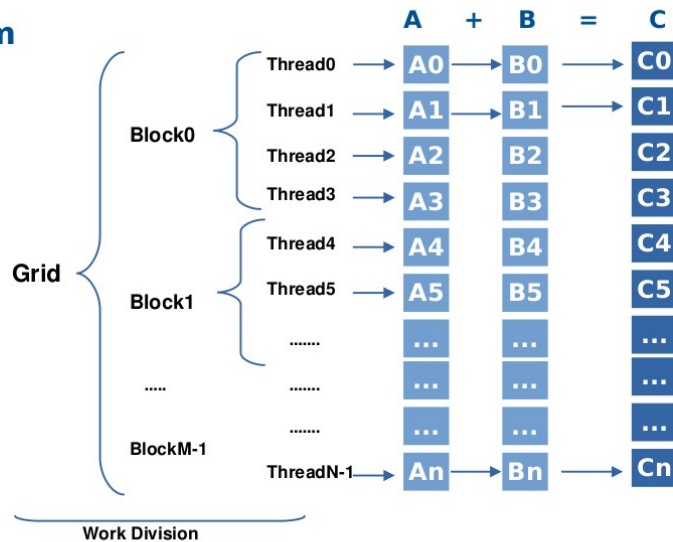
$$C = A + B$$



Lets assume that we want to sum to vectors by utilising paralism of the hardware. We have 2 one dimensional vectors, vector A and vector B and we are going to calculate their sum C.

**Vector Addition Problem**

$$C = A + B$$

**I- Mapping ThreadIndex to DataIndex**One to one,  
Identity**II- Grouping threads**

Typically for parallelisation we need to divide the calculation of vector C, into parallel summations.

That means we need to map threads to the data. In our solution as you can see on the graph a one-to-one mapping, actually an identity function, between thread indices and data indices is used.

Thread0 will sum  $A[0]$  and  $B[0]$  to find  $C[0]$ , thread1 will sum  $A[1]$  and  $B[1]$  to find  $C[1]$  and so on.

The Second issue is how to select or define the grid-block-thread parallelism or in other words determining the alpaka work division. In this representation we have M blocks in grid and each block has 4 threads.

## Vector Addition Code Steps

1. Create Kernel.
2. Decide where will the parallel and non-parallel parts of the code run.
3. Decide how to parallelise (number of blocks and threads).
4. Allocate host and device memory for A,B and C.
5. Copy the memory to the device.
6. Run the kernel
7. Copy the result data back to the host.

From the coding perspective; the vector addition code should have these steps:

1. Create Kernel.
2. Decide where will the parallel and non-parallel parts of the code run.
3. Decide how to parallelise (number of blocks and threads).
4. Allocate host and device memory for A,B and C.
5. Copy the memory to the device.
6. Run the kernel
7. Copy the result data back to the host.

## 1. Define the alpaka Kernel

- Contains the algorithm that is run by each thread
- alpaka Kernels are functors (function-like C++ structs / classes) or lambdas
- Arguments can be pointers and trivially copyable types
- Agnostic to device details

```
// Single header library
#include <alpaka/alpaka.hpp>

#include <iostream>
//! An example kernel: vector addition
class VectorAddKernel
{
public:
    ALPAKA_NO_HOST_ACC_WARNING
    template<typename TAcc, typename TElem, typename TIdx>
    ALPAKA_FN_ACC auto operator()(
        TAcc const& acc, // the accelerator
        TElem const* const A, TElem const* const B, TElem* const C,
        TIdx const& numElements) const -> void
    {
        static_assert(alpaka::Dim<TAcc>::value == 1, "The kernel expects 1-dimensional indices!");
        // Get thread index
        TIdx const gridThreadId(alpaka::getIdx<alpaka::Grid, alpaka::Threads>(acc)[0u]);

        if(gridThreadId < numElements)
        {
            // Use thread index as the data index
            C[gridThreadId] = A[gridThreadId] + B[gridThreadId];
        }
    }
};
```

At the first step kernel is defined. Kernel contains code that is run by each thread. Alpaka kernels are functors, namely structs with specifically implemented function operators or lambdas.

Arguments of the function operator can be pointers or trivially copyable types. You can put many pointers and built-in types in a struct and just pass the struct as a value for example.

Alpaka kernel is agnostic to device details.

As you see, Alpaka is low level and transparent. Abstraction usually associated with being high level but alpaka is low level in that sense. One can access to the thread index directly as in Hip or Cuda.

### Obtaining the indices of threads/blocks inside the Kernel

Grid												
Block 0					Block 1					Block 2		
Threads					Threads					Threads		
0.0	0.1	0.2	0.3	0.4	0.0	0.1	0.2	0.3	0.4	0.0	0.1	0.2
1.0	1.1	1.2	1.3	1.4	1.0	1.1	1.2	1.3	1.4	1.0	1.1	1.2
2.0	2.1	2.2	2.3	2.4	2.0	2.1	2.2	2.3	2.4	2.0	2.1	2.2
3.0	3.1	3.2	3.3	3.4	3.0	3.1	3.2	3.3	3.4	3.0	3.1	3.2

- Index of Thread on the Grid:  

```
auto gridThreadIndex = alpaka::getIdx<alpaka::Grid, alpaka::Threads>(acc);
// gridThreadIndex is {1,9}
```
- Index of Thread on a Block:  

```
auto threadBlockIndex = alpaka::getIdx<alpaka::Block, alpaka::Threads>(acc);
// threadBlockIndex is {1,4}
```
- Index of Block on the Grid:  

```
auto blockGridIndex = alpaka::getIdx<alpaka::Grid, alpaka::Blocks>(acc);
// the blockGridIndex is {1}
```

Obtaining the indices of threads or blocks is easy in alpaka kernel, thanks to usage of function templates.

Using `getIdx` function with different predetermined template arguments would be enough get the thread index in grid or in block. Or the index of block in the grid.

For example the thread [1,4] of the block1 would have a grid index [1,9] because it is on the second row and the 10<sup>th</sup> column in the grid. Note that `alpaka::Grid` and `alpaka::Threads` types are used as template arguments.

Same thread will have a `threadBlock` index [1,4] which shows its coordinates wrt the block.



## 2. Select the Accelerator, Platform and Device

- alpaka provides a number of pre-defined **Accelerators**.

- AccGpuCudaRt for **Nvidia** GPUs
- AccGpuHipRt for **AMD**, **Intel** and **Nvidia** GPUs
- AccGpuSyclIntel for **AMD**, **Intel** and **Nvidia** GPUs

- AccCpuFibers based on Boost.fiber
- AccCpuOmp2Blocks based on OpenMP 2.x
- AccCpuOmp4 based on OpenMP 4.x
- AccCpuTbbBlocks based on TBB
- AccCpuThreads based on std::thread
- AccCpuSycl

- AccFpgaSyclIntel

- Device** instance represents a single physical device

```
auto main() -> int
{
    // Define the index domain.
    using Dim = alpaka::DimInt<1u>; // Dim is set to 1
    using Idx = std::size_t; // Index type is size_t
    // Define the buffer element type
    using Data = std::uint32_t;

    // Define the accelerator
    // It is possible to choose from a set of accelerators:
    // - AccGpuCudaRt, AccGpuHipRt, AccCpuThreads, AccCpuOmp2Threads, AccCpuOmp2Blocks,
    // - AccCpuTbbBlocks, AccCpuSerial
    using Acc = alpaka::AccGpuCudaRt<Dim, Idx>;
    using DevAcc = alpaka::Dev<Acc>;

    // Select a device from platform of Acc
    auto const platform = alpaka::Platform<Acc>{};
    auto const devAcc = alpaka::getDevByIdx(platform, 0);
}
```

Now we can design the code part that will NOT run in parallel, namely the code that will run on the host device which will just trigger the parallel part.

Initially selecting the one of the predefined accelerator types according to your system is needed. There are 3 accelerator types for as the GPU backend, 6 types for the CPU backend and 1 type for the Fpga backend. The accelerator name actually has 2 important parts first one is the device type (CPU or GPU) and the second part the specific API.

// text Accelerator is an abstraction that makes the kernel 'device agnostic'.  
 // easily change the accelerator

Accelerator is a type that is only instantiated on Device not on the host.

Easy management of physical devices

Contains device information

### 3. How to parallelise?

#### I- Get a valid work division from alpaka

Use `getValidWorkDiv` function

- The function divides the full grid-thread extent into blocks.
- Inputs:
  - Full grid-thread extent. (User provides total number of threads needed.)
  - Elements per thread extent
- Most probable `workDiv` in the code will be `{Vec{numElements/1024}, Vec{1024}, Vec{1}}`

#### II - Determine the workdivision manually

- WorkDivision data structure consists 3 vectors:
  - Grid block extent.  
`Vec{numElements/1024}` or `Vec{1,1,numElements/2048}` depending on the number of dimensions.
  - Block thread extent.  
`Vec{1024}` or `Vec{1,1,1024}`
  - Elements per thread is `Vec{1}` or `Vec{1,1,1}`

```
auto const platform = alpaka::PlatformAcc{};
auto const devAcc = alpaka::getDevByIdx(platform, 0);

// Define the work division depending on the data
Idx const numElements(100000);
Idx const elementsPerThread(10);
alpaka::Vec<Dim, Idx> const extent(numElements);

// Let alpaka calculate good block and grid sizes given our full problem extent
alpaka::WorkDivMembers<Dim, Idx> const workDiv = alpaka::getValidWorkDivAcc(*
    devAcc, // device
    extent, // {length, height, depth} of grid. For 1D only length of the vector: {length!
    elementsPerThread,
    false, alpaka::GridBlockExtentSubDivRestrictions::Unrestricted);

// .....

// .....

// Instantiate the kernel function object
VectorAddKernel kernel;
alpaka::execAcc( // Run the kernel execution task
    queue,
    workDiv,
    kernel, alpaka::getPtrNative(bufAccA), alpaka::getPtrNative(bufAccB),
    alpaka::getPtrNative(bufAccC),
    numElements);
// Copy back the result
alpaka::memcpy(queue, bufHostC, bufAccC); // bufHostC includes the result!
alpaka::wait(queue);
}
```

Parallelisation model or work division can be selected using an alpaka function `GetValidWorkDiv`, which takes the full grid-thread extent as argument and divides the given extent into blocks. It takes namely a massive box of threads without any subdivisions and divides this box of threads into blocks. The generated block size is inside the allowed limits which is usually 1024 threads per block for GPUs. Elements per thread extent is a sub index which could be assigned to each thread; if each thread needs a number of additional indices.

Secondly work division can be set by the user. Grid-block, block thread extent and thread-elem extent should be determined by the user.

#### 4. Allocate data vectors A and B on host and device.

- **alpaka::Buf** is multi-dimensional dynamic array.  
It contains
  - *memory*,
  - *size*,
  - the *device* it is located in!
- **alpaka::allocBuf()** allocates memory to the given device.
- **alpaka::View** is used to adapt existing memory, if we already have an STL vector for example we don't need to create Buf, getting a view of existing STL contiguous container would be ok.

```
// Select a device from platform of Acc
auto const platform = alpaka::Platform<Acc>{};
auto const devAcc = alpaka::getDevByIdx(platform, 0);

// Define the work division depending on the data
Idx const numElements(100000);
Idx const elementsPerThread(10);
alpaka::Vec<Dim, Idx> const extent(numElements);
...
// Get the host device for allocating memory on the host.
auto const platformHost = alpaka::PlatformCpu{};
// Get the device directly from CPU platform not from the platform of Acc
auto const devHost = alpaka::getDevByIdx(platformHost, 0);

// Host device type is needed, because it is not known (for the backend it is known in Acc)
using DevHost = alpaka::DevCpu;
// Allocate 3 host memory buffers
using BufHost = alpaka::Buf<DevHost, Data, Dim, Idx>; // Data: uint32_t, Dim:1, Idx:size_t
BufHost bufHostA(alpaka::allocBuf<Data, Idx>(devHost, extent));
BufHost bufHostB(alpaka::allocBuf<Data, Idx>(devHost, extent));
BufHost bufHostC(alpaka::allocBuf<Data, Idx>(devHost, extent));

// Fill host buffers
...

// Allocate 3 buffers on the accelerator
using BufAcc = alpaka::Buf<DevAcc, Data, Dim, Idx>;
BufAcc bufAccA(alpaka::allocBuf<Data, Idx>(devAcc, extent));
BufAcc bufAccB(alpaka::allocBuf<Data, Idx>(devAcc, extent));
BufAcc bufAccC(alpaka::allocBuf<Data, Idx>(devAcc, extent));
```

## The forth step is Allocating memory for data vectors A and B on host and device.

**For allocation of memory alpaka::Buf** is used, it is multi-dimensional dynamic array.

It contains *memory adress*, *size*, the *device* it is the memory allocated!

**alpaka::allocBuf()** allocates memory to the given device.

alpaka::View is used to adapt existing memory, if we already have an STL vector filled with data at host for example we don't need to create Buf, getting a view of existing STL contiguous container would be ok.

### 5.1 Create the Queue for memcpy and kernel task

- alpaka::Queue is "a queue of tasks"
- Queue is always FIFO, everything is sequential inside the queue.
- and more
  - Different queues run in parallel for many devices
  - Used for synchronization
  - Accelerator back-ends can be mixed within a device queue.
- ...

```
// Create a queue on the device, define the synchronization behaviour
alpaka::Queue<Acc, alpaka::Blocking> queue(devAcc);

// Copy from Host to Acc
alpaka::memcpy(queue, bufAccA, bufHostA);
alpaka::memcpy(queue, bufAccB, bufHostB);
alpaka::memcpy(queue, bufAccC, bufHostC);

// Instantiate the kernel function object
VectorAddKernel kernel;
alpaka::exec<Acc>( // Run the kernel execution task
    queue,
    workDiv,
    kernel, alpaka::getPtrNative(bufAccA), alpaka::getPtrNative(bufAccB),
    alpaka::getPtrNative(bufAccC),
    numElements);
// Copy back the result
alpaka::memcpy(queue, bufHostC, bufAccC); // bufHostC includes the result!
alpaka::wait(queue);
```

Create **alpaka::Queue** using the device instance and the accelerator type (e.g GPU)

Queue is similar to stream in Cuda. Alpaka::Queue is always FIFO, everything is sequential inside the queue.

Two queue types: blocking and non-blocking.

Blocking means, the execution of task blocks the caller, in other words when a task is enqueued or executed; the calling thread is blocked. This property of does not affect relation between queues.

*// Blocking-queues block the caller(host) until Device-side command returns.*

*if we create a non blocking queue using a CudaGpu accelerator type and nvidia device and another non blocking queue using the HipGpu accerator type and AMD device; host could execute a task on the first device then without being blocked could execute a task on the second device.*

*Since we are using single queue the operations on queue is always sequential we don't need to think about whether the calling thread is blocked or not*



## 5.2 Copy data vectors to the Device

- **alpaka::memcpy** copies the data from one buffer/view to another buffer or view.
- **alpaka::Buf** knows the device it belongs to.

```
// Allocate 3 buffers on the accelerator
using BufAcc = alpaka::Buf<DevAcc, DataType, Dim, Idx>;
BufAcc bufAccA(alpaka::allocBuf<DataType, Idx>(devAcc, extent));
BufAcc bufAccB(alpaka::allocBuf<DataType, Idx>(devAcc, extent));
BufAcc bufAccC(alpaka::allocBuf<DataType, Idx>(devAcc, extent));

// Create a queue on the device, define the synchronization behaviour
alpaka::Queue<Acc, alpaka::Blocking> queue(devAcc);

// Copy from Host to Acc
alpaka::memcpy(queue, bufAccA, bufHostA);
alpaka::memcpy(queue, bufAccB, bufHostB);
alpaka::memcpy(queue, bufAccC, bufHostC);
```

*because for all operations we used the same queue.*

Copying data vectors to device done by `alpaka::memcpy`. `Memcpy` copies the second buffer argument to first buffer argument. As you would see in the highlighted `memcpy` call; only the buffers allocated to host and device are given to function `memcpy` without stating at which device they are allocated. Because the device information is already inside the buffer data structure.

## 6. Execute the kernel

- Call `alpaka::exec` function
- The result is stored in an `alpaka::Buf`

## 7. Copy result back

- Copy the result in device to the host

```
// Instantiate the kernel function object
VectorAddKernel kernel;

alpaka::exec<Acc>( // Run the kernel execution task
    queue,
    workDiv,
    kernel,
    alpaka::getPtrNative(bufAccA),
    alpaka::getPtrNative(bufAccB),
    alpaka::getPtrNative(bufAccC),
    numElements);
// Copy back the result
alpaka::memcpy(queue, bufHostC, bufAccC); // bufHostC includes the result!
```

On the last 2 steps kernel is executed using the queue and result is copied back to the host. Since everything in the queue is sequential without being blocking and we used the same queue for 2 operations we are sure that there is no problem synchronization.

```
// Single header library
#include <alpaka/alpaka.hpp>
#include <iostream>

// An example kernel: vector addition
class VectorAddKernel
{
public:
    ALPACA_NO_HOST_ACC_WARNING
    template<typename TAcc, typename TElem, typename TIdx>
    ALPACA_FN_ACC auto operator()(
        TAcc const& acc, // the accelerator
        TElem const& const A,
        TElem const& const B,
        TElem const C,
        TIdx const& numElements) const -> void
    {
        static_assert(alpaka::Dim(TAcc)::value == 1, "Kernel expects 1-dimensional indices!");
        // Get thread index
        TIdx const gridThreadIdx(alpaka::getIdx<alpaka::Grid, alpaka::Threads>(acc)[0]);

        if(gridThreadIdx < numElements)
        {
            // Use thread index as the data index
            C[gridThreadIdx] = A[gridThreadIdx] + B[gridThreadIdx];
        }
    }
};

auto main() -> int
{
    using Dim = alpaka::DimInt<1>; // Define the index domain
    using Idx = std::size_t; // Index type of the threads and buffers
    using DataType = std::uint32_t; // Define the buffer element type

    // Define the accelerator: AccCpuCudaRt, AccCpuHlRt,
    // AccCpuThreads, AccCpuDmp2Threads, AccCpuDmp2Blocks, AccCpu10bBlocks AccCpuSerial
    using Acc = alpaka::AccCpuCudaRt<Dim, Idx>;
    using DevAcc = alpaka::DevAcc<Acc>;

    // Select a device from platform of Acc
    auto const platform = alpaka::Platform<Acc>{};
    auto const devAcc = alpaka::getDevByIdx(platform, 0);

    // Define the work division depending on the data
    Idx const numElements(100000);
    Idx const elementsPerThread(10);
    alpaka::Vec<Dim, Idx> const extent(numElements);

    // Let alpaka calculate good block and grid sizes given our full problem extent
    alpaka::WorkDivMembers<Dim, Idx> const workDiv = alpaka::getValidWorkDiv<Acc>(*
        devAcc, // device
        extent, // (length, height, depth) of grid. For 1D only length of the vector!
        elementsPerThread,
        false, alpaka::GridBlockExtentSubDivRestrictions::Unrestricted);

    // Get the host device for allocating memory on the host.
    auto const platformHost = alpaka::PlatformCpu{};
    // Get the device directly from CPU platform not from the platform of Acc
    auto const devHost = alpaka::getDevByIdx(platformHost, 0);

    // Host device type is needed, because it is not known (for the backend it is known in Acc)
    using DevHost = alpaka::DevCpu;
    // Allocate 3 host memory buffers
    using BufHost = alpaka::Buf<DevHost, DataType, Dim, Idx>;
    BufHost bufHostA(alpaka::allocBuf<DataType, Idx>(devHost, extent));
    BufHost bufHostB(alpaka::allocBuf<DataType, Idx>(devHost, extent));
    BufHost bufHostC(alpaka::allocBuf<DataType, Idx>(devHost, extent));

    // Fill the buffers
    for(Idx i{0}; i < numElements; ++i)
    {
        bufHostA[i] = random(); bufHostB[i] = random(); bufHostC[i] = 0;
    }

    // Allocate 3 buffers on the accelerator
    using BufAcc = alpaka::Buf<DevAcc, DataType, Dim, Idx>;
    BufAcc bufAccA(alpaka::allocBuf<DataType, Idx>(devAcc, extent));
    BufAcc bufAccB(alpaka::allocBuf<DataType, Idx>(devAcc, extent));
    BufAcc bufAccC(alpaka::allocBuf<DataType, Idx>(devAcc, extent));

    // Create a queue on the device, define the synchronization behaviour
    alpaka::Queue<Acc, alpaka::Blocking> queue(devAcc);

    // Copy from Host to Acc
    alpaka::memcpy(queue, bufAccA, bufHostA);
    alpaka::memcpy(queue, bufAccB, bufHostB);
    alpaka::memcpy(queue, bufAccC, bufHostC);

    // Instantiate the kernel function object
    VectorAddKernel kernel;
    alpaka::exec<Acc>(& // Run the kernel execution task
        queue,
        workDiv,
        kernel, alpaka::getPtrNative(bufAccA), alpaka::getPtrNative(bufAccB),
        alpaka::getPtrNative(bufAccC),
        numElements);

    // Copy back the result
    alpaka::memcpy(queue, bufHostC, bufAccC); // bufHostC includes the result!
}
```



## Parallel vector addition code

```
// Single header library
#include <alpaka/alpaka.hpp>

#include <iostream>

// An example kernel: vector addition
class VectorAddKernel
{
public:
    ALPACA_NO_HOST_ACC_WARNING
    template<typename TAcc, typename TElem, typename TIdx>
    ALPACA_FN_ACC auto operator()(
        TAcc const& acc, // the accelerator
        TElem const& A,
        TElem const& B,
        TIdx const& numElements) const -> void
    {
        static_assert(alpaka::DimTAcc::value == 1, "Kernel expects 1-dimensional indices");
        // Get thread index
        TIdx const gridThreadIdx(alpaka::getIdx<alpaka::Grid, alpaka::Threads>(acc)[0]);

        if(gridThreadIdx < numElements)
        {
            // Use thread index as the data index
            C[gridThreadIdx] = A[gridThreadIdx] + B[gridThreadIdx];
        }
    }
};

auto main() -> int
{
    using Dim = alpaka::DimInt<1>; // Define the index domain
    using Idx = std::size_t; // Index type of the threads and buffers
    using DataType = std::uint32_t; // Define the buffer element type

    // Define the accelerator: AccGpuCudaRt, AccGpuPrt,
    // AccCpuThreads, AccCpuOmp2Threads, AccCpuOmp2Blocks, AccCpuOmp2Blocks, AccCpuOmp2Blocks
    using Acc = alpaka::AccGpuCudaRt<Dim, Idx>;
    using DevAcc = alpaka::Dev<Acc>;

    // Select a device from platform of Acc
    auto const platform = alpaka::Platform<Acc>{};
    auto const devAcc = alpaka::getDevByIdx(platform, 0);

    // Define the work division depending on the data
    Idx const numElements(1000000);
    Idx const elementsPerThread(1000);
    alpaka::Vec<Dim, Idx> const extent(numElements);
```

Single  
header

kernel

Select  
accelerator  
and the  
device (GPU)

```
// Let alpaka calculate good block and grid sizes given our full problem extent
alpaka::WorkDivMembers<Dim, Idx> const workDiv = alpaka::getValidWorkDiv<Acc>(
    devAcc, // device
    extent, // (length, height, depth) of grid. For 1D only length of the vector
    alpaka::getPerThread(), // threads per thread
    false, alpaka::GridBlockExtentSubDivRestrictions::Unrestricted);

// Get the host device for allocating memory on the host.
auto const platformHost = alpaka::PlatformCpu{};
// Get the device directly from CPU platform not from the platform of Acc
auto const devHost = alpaka::getDevByIdx(platformHost, 0);

// Host device type is needed, because it is not known (for the backend it is known in Acc)
using DevHost = alpaka::DevCpu;
// Allocate 2 host memory buffers
using BufHost = alpaka::Buf<DevHost, DataType, Dim, Idx>;
BufHost bufHostA(alpaka::allocBuf<DataType, Idx>(devHost, extent));
BufHost bufHostB(alpaka::allocBuf<DataType, Idx>(devHost, extent));
BufHost bufHostC(alpaka::allocBuf<DataType, Idx>(devHost, extent));

// Fill the buffers
for(Idx i(0); i < numElements; ++i)
{
    bufHostA[i] = randomA; bufHostB[i] = randomB; bufHostC[i] = 0;
}

// Allocate 2 buffers on the accelerator
using BufAcc = alpaka::Buf<DevAcc, DataType, Dim, Idx>;
BufAcc bufAccA(alpaka::allocBuf<DataType, Idx>(devAcc, extent));
BufAcc bufAccB(alpaka::allocBuf<DataType, Idx>(devAcc, extent));
BufAcc bufAccC(alpaka::allocBuf<DataType, Idx>(devAcc, extent));

// Create a queue on the device, define the synchronization behaviour
alpaka::Queue<Acc, alpaka::Blocking, queue(devAcc)> queue;

// Copy from host to Acc
alpaka::memcpy(queue, bufHostA, bufHostA);
alpaka::memcpy(queue, bufHostB, bufHostB);
alpaka::memcpy(queue, bufHostC, bufHostC);

// Instantiate the kernel function object
VectorAddKernel kernel;
alpaka::exec<Acc>( // Run the kernel execution task
    queue,
    workDiv,
    kernel, alpaka::getPtrNative(bufHostA), alpaka::getPtrNative(bufHostB),
    alpaka::getPtrNative(bufHostC),
    numElements);

// Copy back the result
alpaka::memcpy(queue, bufHostC, bufAccC); // bufHostC includes the result
```

Get Work  
division

Get host  
device (CPU)

allocate  
memory  
at host (CPU)

allocate  
memory at  
device (GPU)

Copy vectors  
to device  
(GPU)

Execute  
kernel

Copy  
result  
to host (CPU)




### Programing Tips

- You can do printf debugging; but can **not use** `std::cout` in alpaka Kernel
- If you want to pass multi-dimensional data to kernel, use `mdspan` (enable it via cmake option).  
(If you don't use `mdspan`; you will need to take care of alignment/pitch values. Pass the pointer, extents and the pitch.)
- A kernel can be run directly by `exec` function or can be **enqueued** as a task.
- If there are unused number of dimensions in workdiv; use 1, for that dimension.  
`auto blockThreadExtent = alpaka::Vec<TDim3D,Idx>{1u,1u,128u};`
- Vendor specific profiling and debugging tools can be used directly on compiled alpaka (e.g. `nsys`, `rocprowf` ...)
- If you debug GPU code try to compile your code for CPU; and use CPU debugger tools  
(Change acc type to CPU accelerators then debug using `gdb` and similar tools.)


I'd like to give some programming tips for users of Alpaka.

## alpaka Performance

Mean GPU and CPU utilization on NVIDIA A40 GPU



Peak memory usage on NVIDIA A40 GPU


**Source: Evaluating Performance Portability with the CMS Heterogeneous Pixel Reconstruction code**

N. Andriotis<sup>1</sup>, A. Bocci<sup>2</sup>, E. Cano<sup>3</sup>, L. Cappelli<sup>3</sup>, M. Dewing<sup>1</sup>, T. Di Pilato<sup>4,5</sup>, J. Esseiva<sup>1</sup>, L. Ferragina<sup>6</sup>, G. Hugo<sup>2</sup>, M. Kortelainen<sup>3</sup>, M. Kwok<sup>3</sup>, J. J. Olivera Loyola<sup>10</sup>, F. Pantaleo<sup>2</sup>, A. Perego<sup>11</sup>, W. Redjet<sup>2,12</sup>  
<sup>1</sup>BSC <sup>2</sup>CERN <sup>3</sup>INFN Bologna <sup>4</sup>ANL <sup>5</sup>CASUS <sup>6</sup>University of Geneva <sup>7</sup>LBNL <sup>8</sup>University of Bologna <sup>9</sup>FNAL <sup>10</sup>ITESM <sup>11</sup>University of Milano Bicocca <sup>12</sup>RWTH AACHEN

CHEP 2023

[https://indico.jlab.org/event/459/contributions/11824/attachments/9281/14171/20230511-CHEAP23\\_CMSPortability.pdf](https://indico.jlab.org/event/459/contributions/11824/attachments/9281/14171/20230511-CHEAP23_CMSPortability.pdf)

A study for understanding the performance of Alpaka was carried out by alpaka contributors and maintainers at CERN. 3 parallel programming tools are compared: Alpaka, a similar abstraction tool called Kokkos and Cuda itself.

The 2 graphs on the lefts show the GPU and CPU utilisation. In the case of mean GPU utilisation; performances of 3 tools were similar as you can see on the left most graph. But the CPU utilisation of alpaka was much better.

The 2 graphs on the right show the peak memory usage. And the memory usage of Alpaka was much better for the GPU and slightly better for the CPU compared to Kokkos.

This performance analysis is done by CERN in 2021.



## alpaka in the wild...

**PICongGPU:**

- Fully relativistic, manycore, 3D3V particle-in-cell (PIC) code
- Implements central algorithms in plasma physics
- Scalable to more than 18,000 GPUs
- Developed at Helmholtz-Zentrum Dresden-Rossendorf



<https://github.com/ComputationalRadiationPhysics/picongpu>



PicongGPU project is one of the most important users of alpaka. It is "particle in cell" code and mainly does laser plasma acceleration.

It is run on some of the fastest super computers.

The graphs show how the performance of picongpu scales as the number of nodes grow according to tests performed on LUMI super computer. In the first graph the problem size increased with the number of size in the second one or the one below the problem size is fixed. Of course there is an underutilization of capacity.

Hence Alpaka makes PicOnGPU to perform well on various different HPC platforms.

### How to start using alpaka

- Don't write code initially on **cuda** because **alpaka** is already low level!
- Use **alpaka directly** by using examples and the cheat-sheet.
- **BUT if you already have a codebase in cuda**, converting to **cupla** can be a fast solution to benefit from alpaka features! Cupla is a member of alpaka group of softwares.

**cupla** - C++ User interface for the Platform Independent Library



<https://github.com/alpaka-group/cupla>

## Cuda to portable C++ code

- Change the suffix \*.cu of the CUDA source files to \*.cpp
- Remove `#include <cuda_runtime.h>` and other cuda specific include files.
- Add `#include <cuda_to_cupla.hpp>`

### Cuda

```

Kernel Function
template<int blockSize>
__global__ void fooKernel(int * ptr, float value)
{ // ... }

Kernel call at host
dim3 gridSize(42,1,1);
dim3 blockSize(256,1,1);
fooKernel<16><<< gridSize, blockSize, 0, 0>>>(ptr, 23);

Device function
template<typename TElem>
__device__ int deviceFunction(TElem x)
{ // ... }
// call
auto result = deviceFunction(x);

Shared memory
__shared__ int foo;
__shared__ int fooCArray2D[4][32];
    
```

### Cupla

```

Kernel Functor
template<int blockSize>
struct fooKernel {
    template<typename TAcc>
    ALPAKA_FN_ACC void operator()(TAcc const & acc, int *
    const ptr, float const value) const
    { // ... }
};

Kernel call at host
dim3 gridSize(42,1,1); dim3 blockSize(256,1,1);
CUPLA_KERNEL(fooKernel<16>)(gridSize, blockSize, 0, 0)(ptr,
23);

Device function
template< typename TAcc, typename TElem >
ALPAKA_FN_ACC int deviceFunction(TAcc const & acc, TElem
x)
{ // ... }
// call
auto result = deviceFunction(acc, x);

Shared memory
sharedMem(foo, int);
sharedMem(fooCArray2D, cupla::Array< cupla::Array<int,4>, 32>);
    
```

## For converting Cuda to portable C++ code

Initially user needs to change the suffix \*.cu of the CUDA source files to \*.cpp


Remove the **#include <cuda\_runtime.h> statement** and other include statements for the cuda specific include files.

And then the user has to include the file **cuda\_to\_cupla.hpp**

As you can see cuda and cupla code is quite similar. Cuda kernel function is converted to a kernel functor. Kernel call in cupla needs CUPLA\_KERNEL macro and as an example shared memory declaration need sharedMem function. A 2d array in shared memory is achieved by a cupla array of items of type cupla array.

## Community and Long Term Support

- Partners using and contributing to alpaka



- alpaka is a part of Helmholtz Roadmap 2027-2034

Before finishing I would like list the users of Alpaka library. CERN is a very important user and contributor, DLR is using alpaka but their codes are not on public domain. HZDR is an important user by creating PicOnGPU and HZDR is also directly contributing to develop and maintain alpaka. And lastly Helmholtz Zentrum Berlin has recently started using alpaka.

On the other hand since Alpaka is a part of Helmholtz Roadmap from 2027 to 2034, a long term support is already secured for Alpaka.

```
// examples select which backend
// CERN code links
// mdspan queue examples
```


### As a summary

#### Without alpaka

- Multiple hardware types are available from different vendors (CPUs, GPUs, ...)
- Increasingly heterogeneous hardware configurations available
- Platforms not inter-operable → parallel programs not easily portable

#### alpaka: one API to rule them all

- **Abstraction** (not hiding!) of the underlying hardware & software platforms
  - AMD, Nvidia, Intel GPUs, Different CPU parallelisations like TbbBlocks, OpenMP, Threads
- **Easy change of the backend in Code**
- **Built down to the same machine code with the vendor solutions**
- **Zero abstraction overhead for Kernel execution!**
- **Heterogenous Programming:** Using different backends in a synchronized manner



As we mentioned before currently HPC Platforms are not interoperable, or in other words programs are not portable.

**Alpaka provides one API to support all different GPUs and CPU backends.**

**Alpaka provides Abstraction** (but not hiding!) of the underlying hardware, compiler and OS

Alpaka does not have default device, built-in functions, language extensions, default stream like in cuda

**It is Easy to change the backend in code**

**Alpaka code use directly of vendor APIs, does not depend on “unified APIs”. Produces the same code that a vendor API would generate. It is not emulating. For example alpaka user can use vendor profilers and debuggers (Cuda,HIP...) for his alpaka code!**

Supported GPU Backends are Hip (AMD), Cuda (NVidia), SYCL (Intel GPUs)

CPU Backends: OpenMp, Threads, TbbBlocks

## Zero abstraction overhead for Kernel execution!

## Heterogenous Programming: Using different backends in a synchronized manner



### If you use alpaka for your research, please cite one of the following publications:

Matthes A., Widera R., Zenker E., Worpitz B., Huebl A., Bussmann M. (2017): Tuning and Optimization for a Variety of Many-Core Architectures Without Changing a Single Line of Implementation Code Using the alpaka Library. In: Kunkel J., Yokota R., Tauber M., Shalf J. (eds) High Performance Computing. ISC High Performance 2017. Lecture Notes in Computer Science, vol 10524. Springer, Cham, DOI: [10.1007/978-3-319-67630-2\\_36](https://doi.org/10.1007/978-3-319-67630-2_36).

E. Zenker et al., "alpaka – An Abstraction Library for Parallel Kernel Acceleration", 2016 IEEE International Parallel and Distributed Processing Symposium Workshops (IPDPSW), Chicago, IL, 2016, pp. 631 – 640, DOI: [10.1109/IPDPSW.2016.50](https://doi.org/10.1109/IPDPSW.2016.50).

Worpitz, B. (2015, September 28). Investigating performance portability of a highly scalable particle-in-cell simulation code on various multi-core architectures. Zenodo. DOI: [10.5281/zenodo.49768](https://doi.org/10.5281/zenodo.49768).

**Thank you!**

**You can contact us for any of your requests or questions about alpaka!**



If you use alpaka for your research please cite one of the publications.

Thank you for your attention. And Please feel free to contact us FOR any of your requests or questions about alpaka. It doesn't matter if it is an installation issue, a bug or a performance problem.







































